



BigQuery Deep Dive

H&M - May 2025

Henrik Warfvinge - Data & AI Architect
henrikw@google.com

Google Cloud

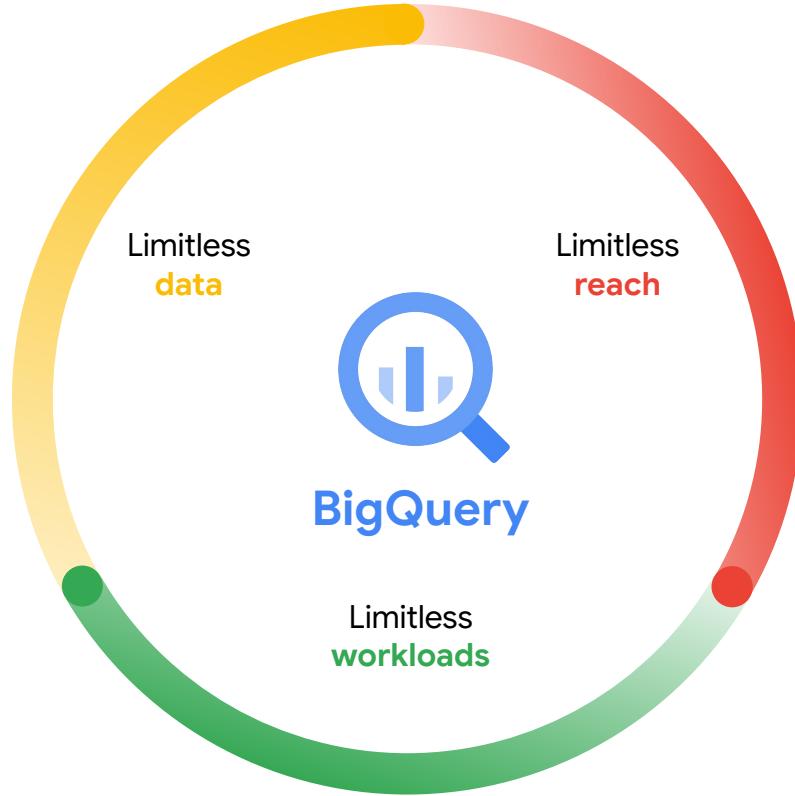


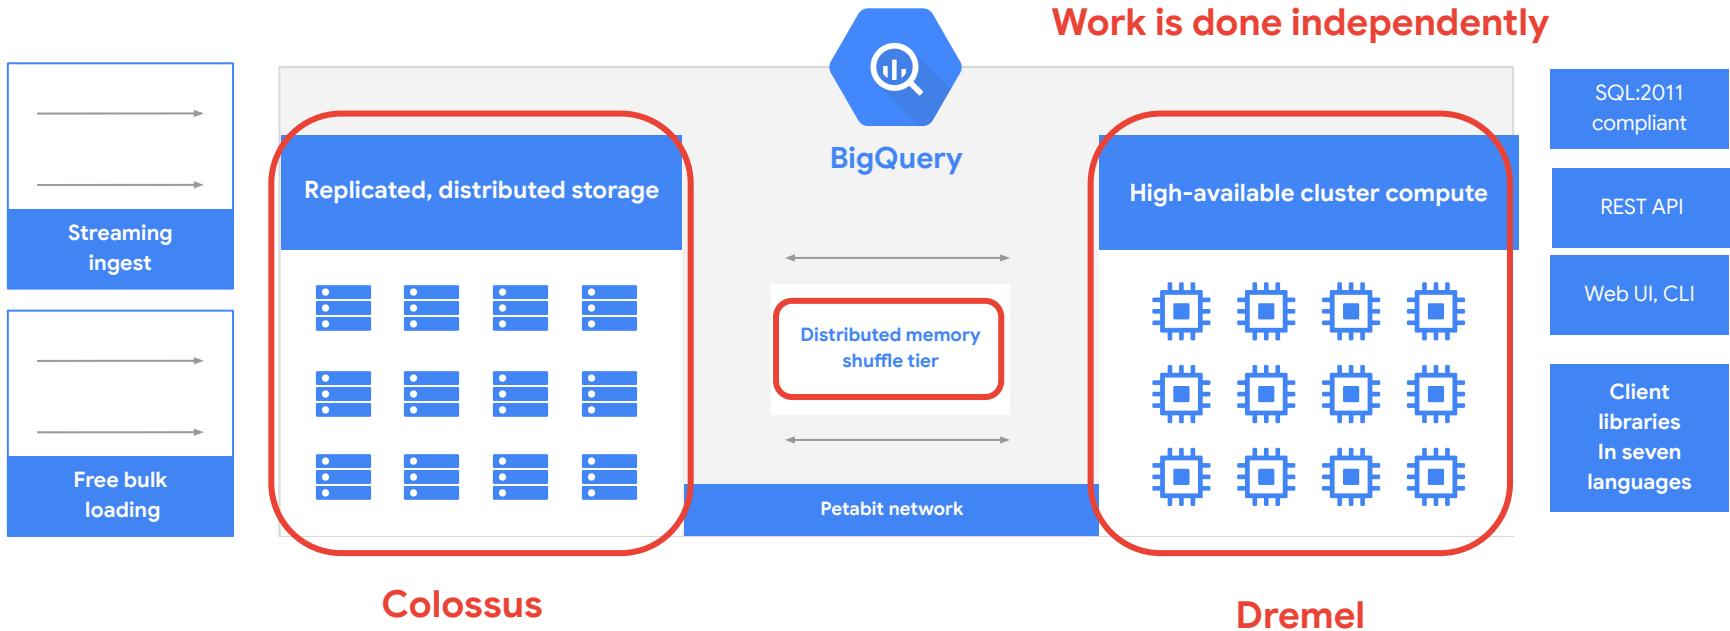


Recap

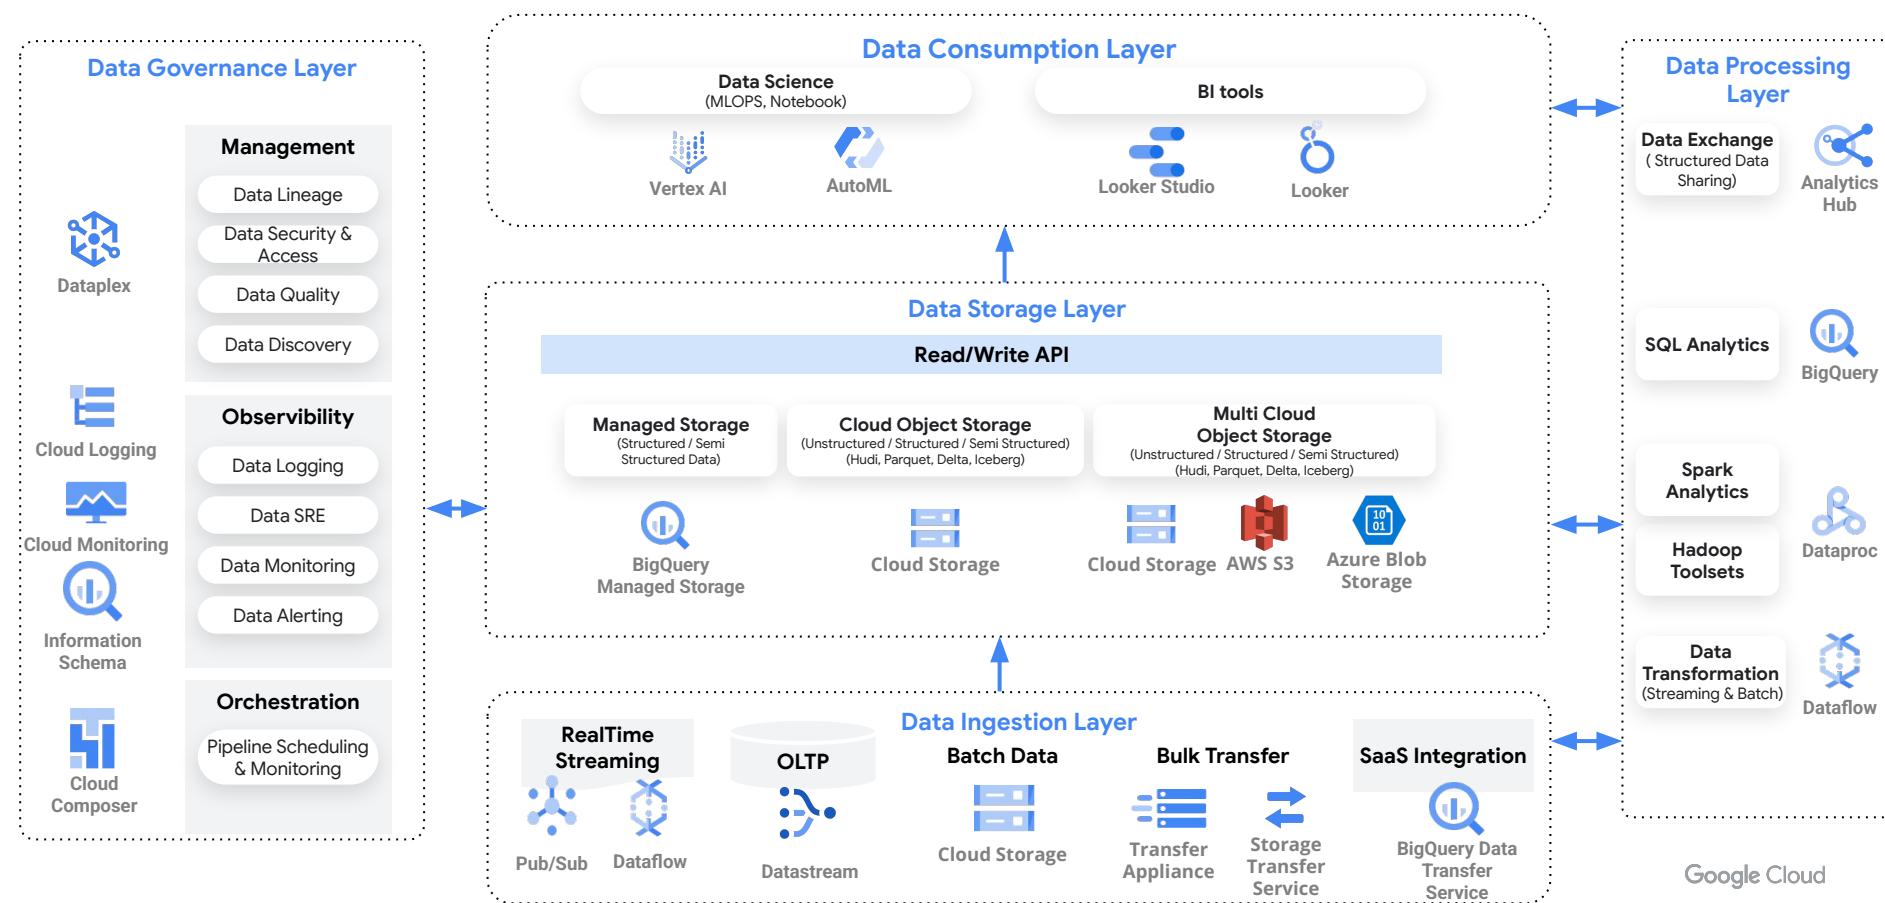
BigQuery

The core of Google's
Data Cloud to power
your **data-driven**
innovation.





Google Analytics Lakehouse Architecture



Agenda

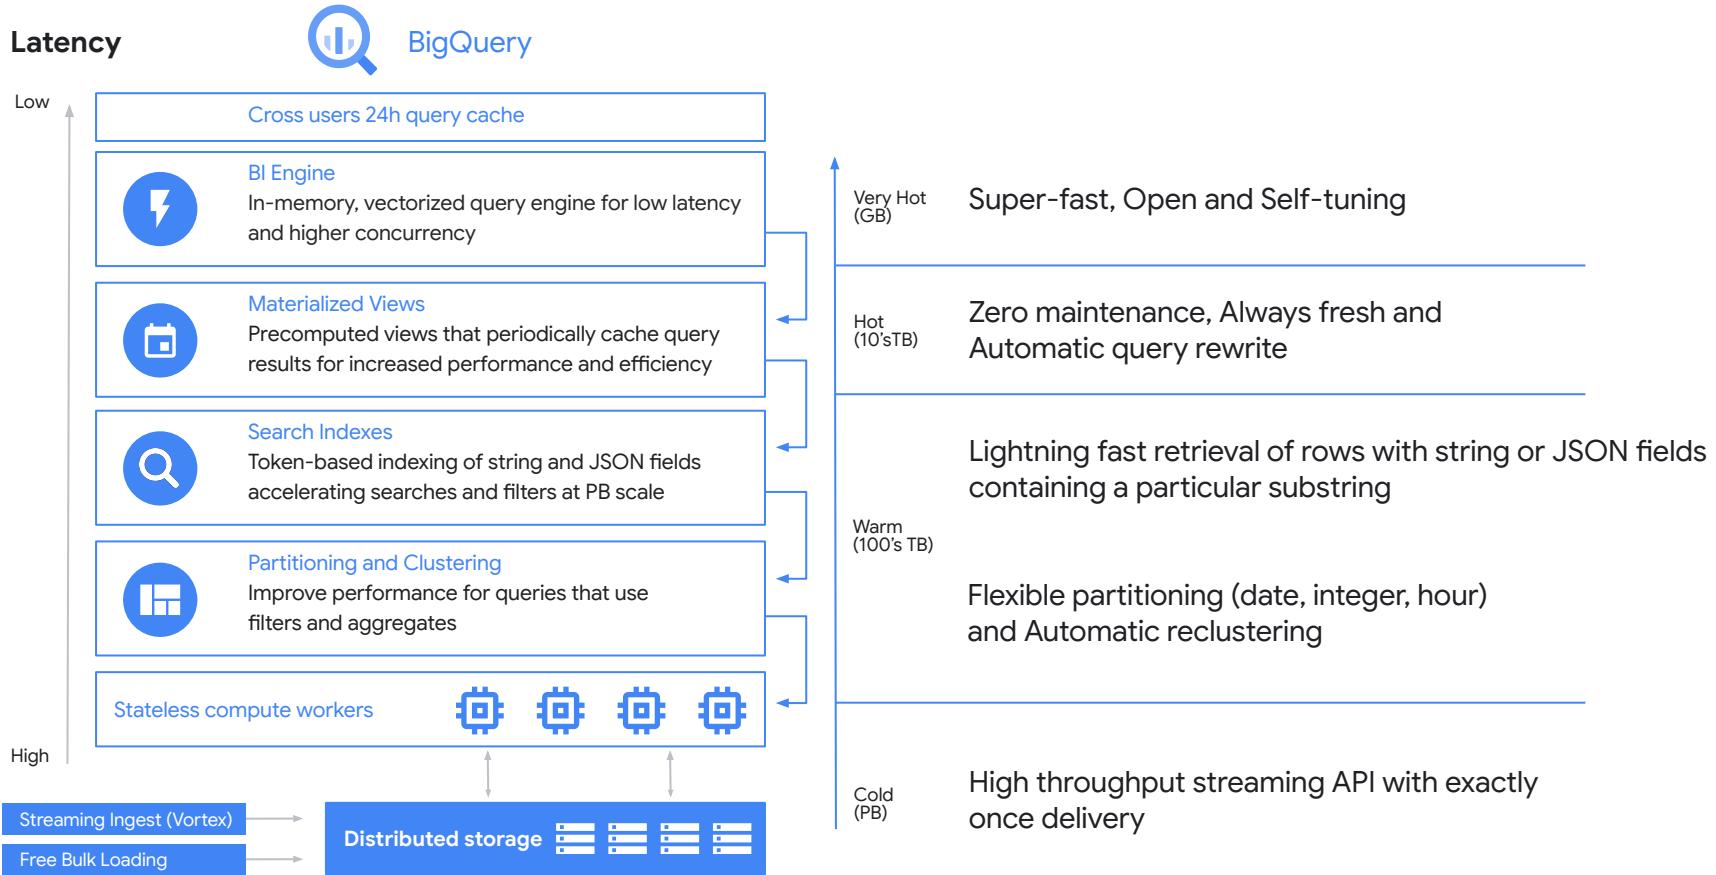
- Storage
 - Partitioning
 - Clustering
 - Search indexes
 - Materialized views
 - Bi Engine
- Compute
 - Slots
 - Query execution
- SQL best practices
- Q&A



Going Deep

Storage Internals and optimizations

Adaptive Caching layers in BigQuery



Key features of BigQuery Storage

Compressed

Industry leading compression, result of over a decade of innovation in storage optimization technology including proprietary columnar compression, automatic data sorting, clustering and compaction.

Encrypted

BigQuery automatically encrypts all data before it is written to disk. You can provide your own encryption key or let Google manage the encryption key.

Managed

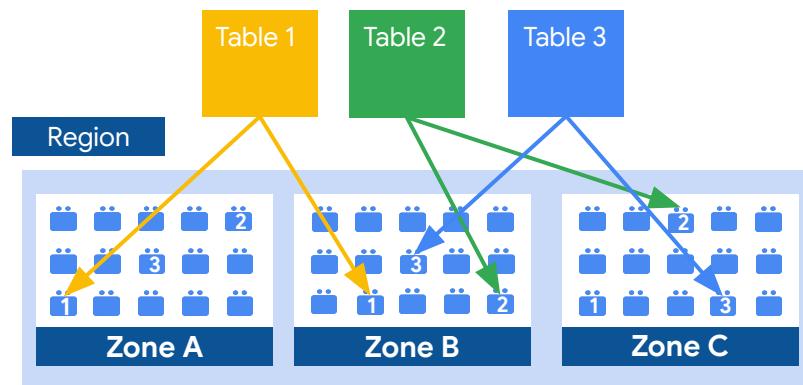
As a completely managed service, you don't need to provision storage resources or reserve units. BigQuery automatically allocates storage when you load data into the system. You only pay for the amount of storage that you use.

Efficient

BigQuery storage uses an innovative adaptive file sizing and efficient columnar format that is optimized for analytic workloads.

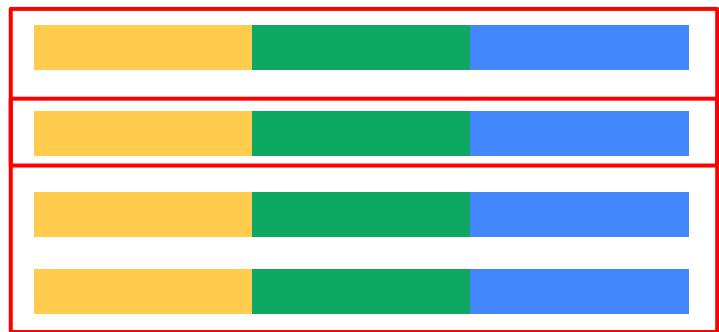
Durable

BigQuery storage is designed for 99.999999999% (11 9's) annual durability. It replicates your data across multiple availability zones to protect from data loss due to machine-level or zonal failures.



Colossus (distributed file system)

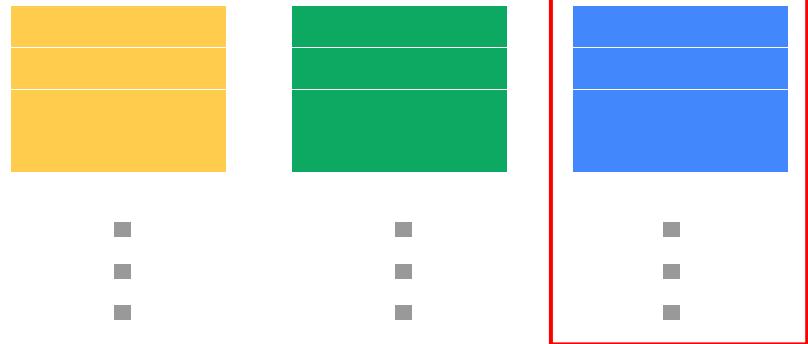
Traditional RDBMS storage



Record-oriented storage

OLTP

BigQuery storage



Columnar storage

OLAP

Data Backup and Recovery

Making it easy to recover from mistakes

GA

Time Travel

Configurable 2-7 day window to access all deleted or changed data.

GA

Table snapshots

Read-only copies of base tables that can be restored for modifications.
Useful for logical backups or archiving beyond time-travel window.

GA

Table clones

Read/write copies of base tables. Useful for testing with production databases.

Preview

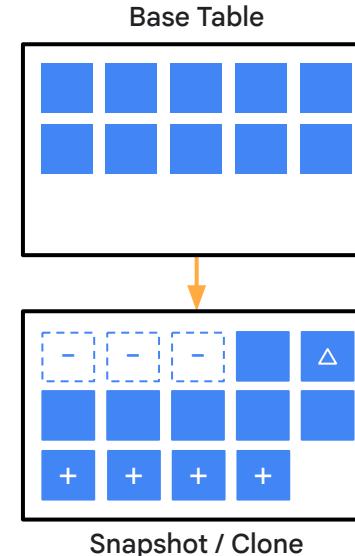
Cross-region table copy

Copies entire table (including CMEK) between regions

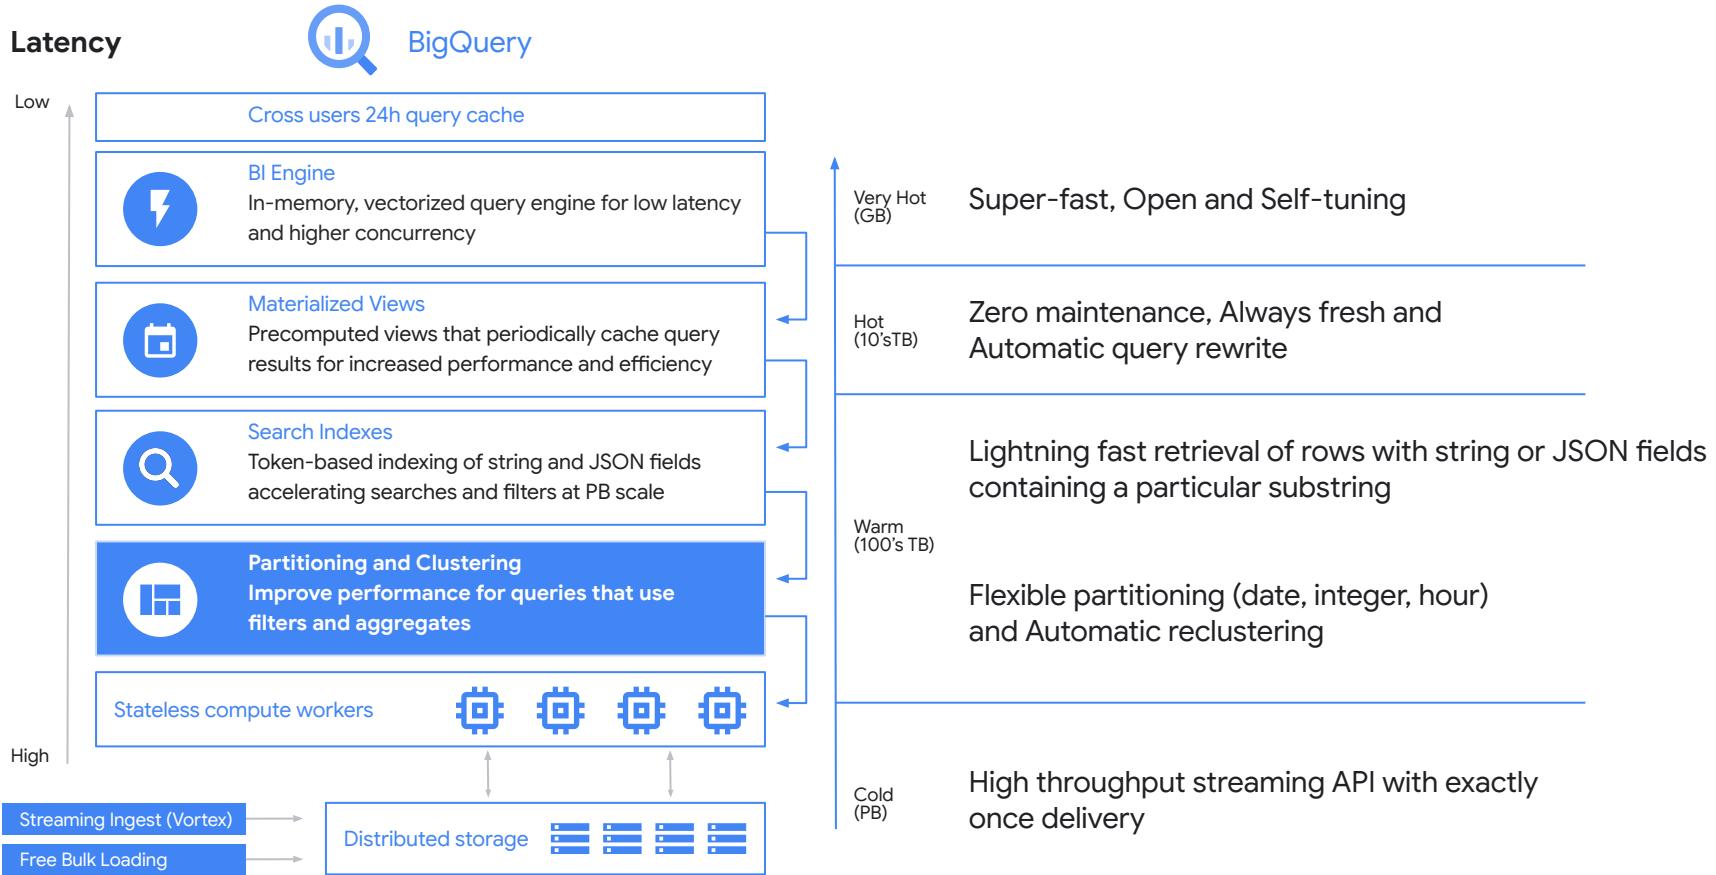
Preview

Dataset undelete

Easily recover entire datasets (all objects) within time travel window.



Adaptive Caching layers in BigQuery

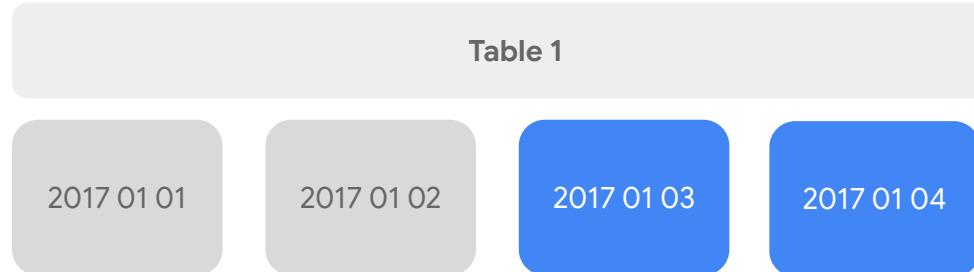


Storage Optimization

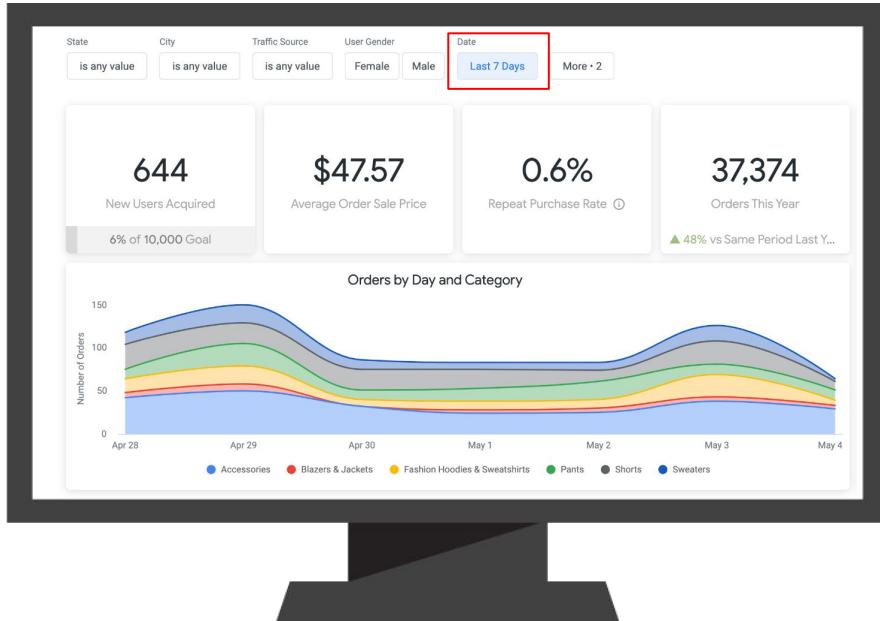
Partitioning

- Ingestion Time
- Date / Timestamp Column
- Integer Column

Table 1



```
SELECT ...
WHERE date >= "20170103"
```



CREATE TABLE

```
project-id.my_dataset.order_items  
(  
    ORDER_ID INT64,  
    ...  
)  
PARTITION BY DATE(created_at)
```

Storage Optimization

Clustering



Customer_ID	Created_At	Order_ID
1	2017-01-01	1
1	2017-01-01	1
1	2017-01-02	2
2	2017-01-01	3
3	2017-01-02	4
3	2017-01-02	4

Table 1

Customers
1-2

Customers
3-4

Customers
5-6

SELECT ...
WHERE customer_id = 1

Partitioning vs. Clustering

Partitioning



Lower Cardinality, Lots of Data

Clustering



Higher Cardinality, Doesn't Need a Lot of Data

Table 1

2017 01 01

2017 01 02

2017 01 03

2017 01 04

Table 1

Customers
1-2

Customers
3-4

Customers
5-6

**SELECT ...
WHERE date >= "20170103"**

**SELECT ...
WHERE customer_id = 1**

(Max 10,000 partitions per table)

Partitioning and Clustering in the same query

```
CREATE TABLE `project-id.my_dataset.order_items`  
(  
    ORDER_ID INT64,  
    ...  
)  
PARTITION BY DATE(created_at)  
CLUSTER BY customer_id
```

Table 1

2017 01 01

2017 01 02

2017 01 03

2017 01 04

SELECT ...
WHERE
date >= "20170103"
AND customer_id = 1

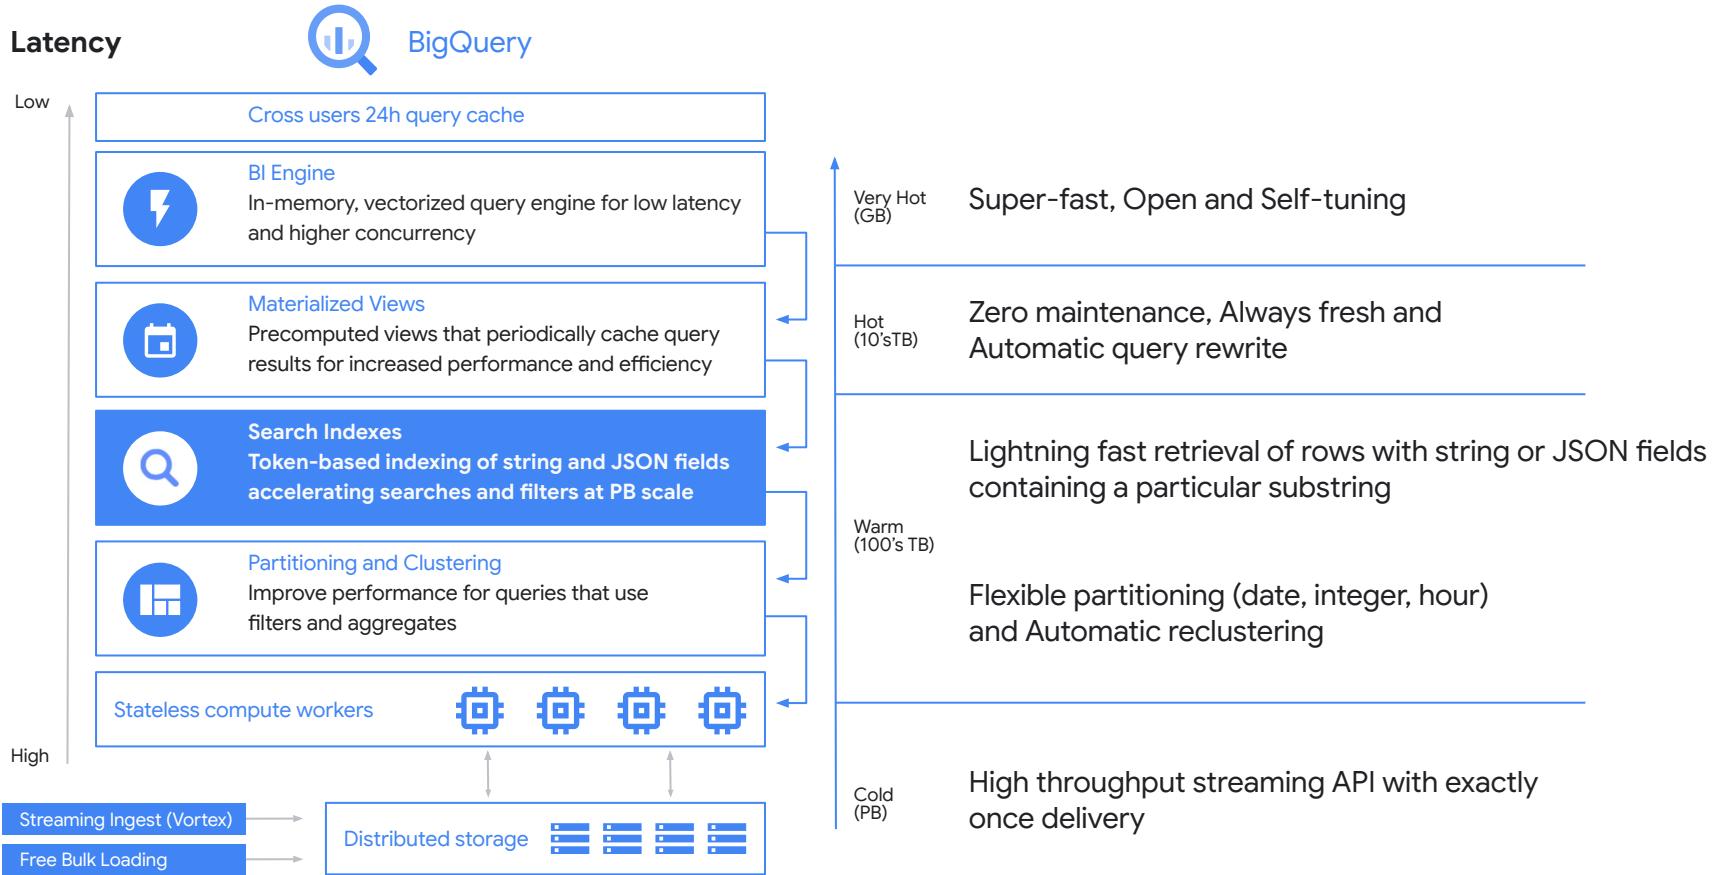
Customers
1-2

Customers
3-4

Customers
1-2

Customers
3-4

Adaptive Caching layers in BigQuery



Search Indexes for BigQuery



```
CREATE SEARCH INDEX bq_demo_index  
ON bq_table
```

SQL

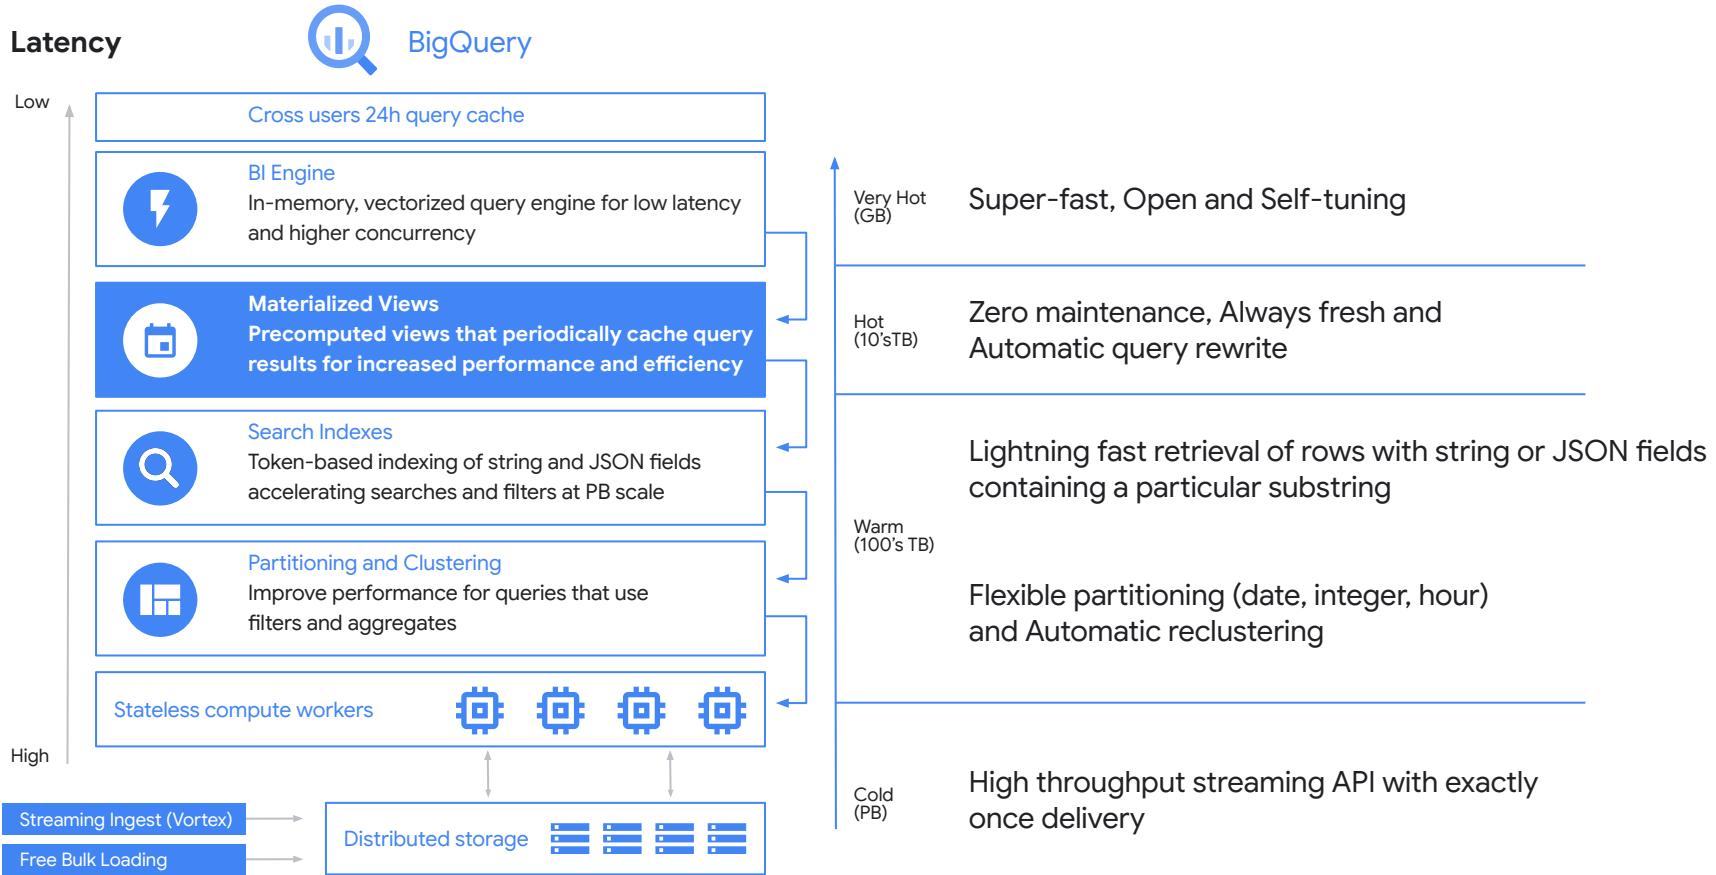
```
SELECT * FROM bq_table t  
WHERE SEARCH (t, 'term')  
    t.field = 'term'  
    t.field like 'term'  
    t.field IN('term1', 'term2')  
    ...
```

Use Search Indexes to accelerate performance of “needle in the haystack” queries (point lookups)

Indexes auto-refresh as data changes and you are always guaranteed correct results. No management required.

Business Case	Example
Queries with filters on JSON	Select only rows where one of the JSON's sub-elements contains a value
Dashboards that require highly selective filters	Slice and dice data by different inventory items
Identify targeted subsets of data within a large dataset	Create a cohort of patients with a precise gene mutation.
Regulatory processes that require finding specific data elements	Find and delete a user's records for GDPR
Developer troubleshooting	Identify the application log entries associated with a particular error code
Security audits	Review all network activity from a specific IP address

Adaptive Caching layers in BigQuery



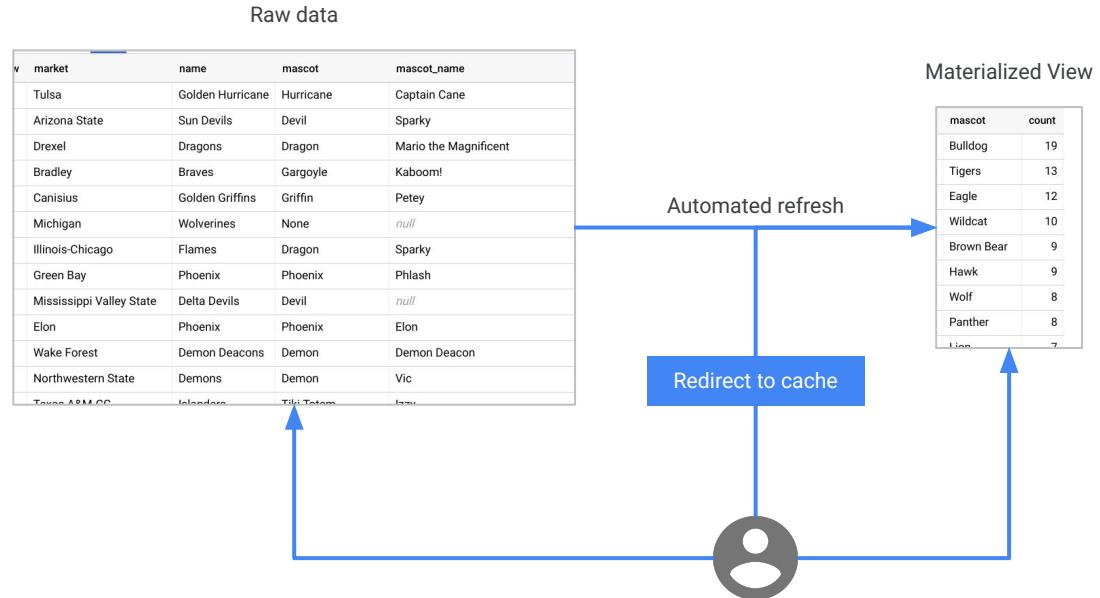
BigQuery Materialized Views

Performance - MVs are a caching layer for faster, cheaper queries

Ease of use - users define MVs, BigQuery automatically maintains and refreshes MVs

Power - Queries are automatically redirected to MVs when applicable

Consistency - data is never stale



Materialized Views - Query Rewrite

Base_Table				
<u>Col 1</u>	<u>Col 2</u>	<u>Col 3</u>	<u>Col 4</u>	<u>Col 5</u>

BQ
Optimizer

```
SELECT  
COL1, SUM(COL3)  
FROM Base_Table  
GROUP BY COL1
```

Materialized_View		
<u>Col 1</u>	<u>Col 3</u>	<u>Col 4</u>

Run Query

Incremental Materialized Views

SQL Limitations

Unsupported SQL features

- Left/right/full outer joins.
- Self-joins (joins using the same table more than once).
- Window functions.
- ARRAY subqueries.
- UNION_ALL
- Non-deterministic functions such as RAND(), CURRENT_DATE(), SESSION_USER(), or CURRENT_TIME().
- User-defined functions (UDFs).
- TABLESAMPLE.
- FOR SYSTEM_TIME AS OF.

Non Incremental Materialized Views

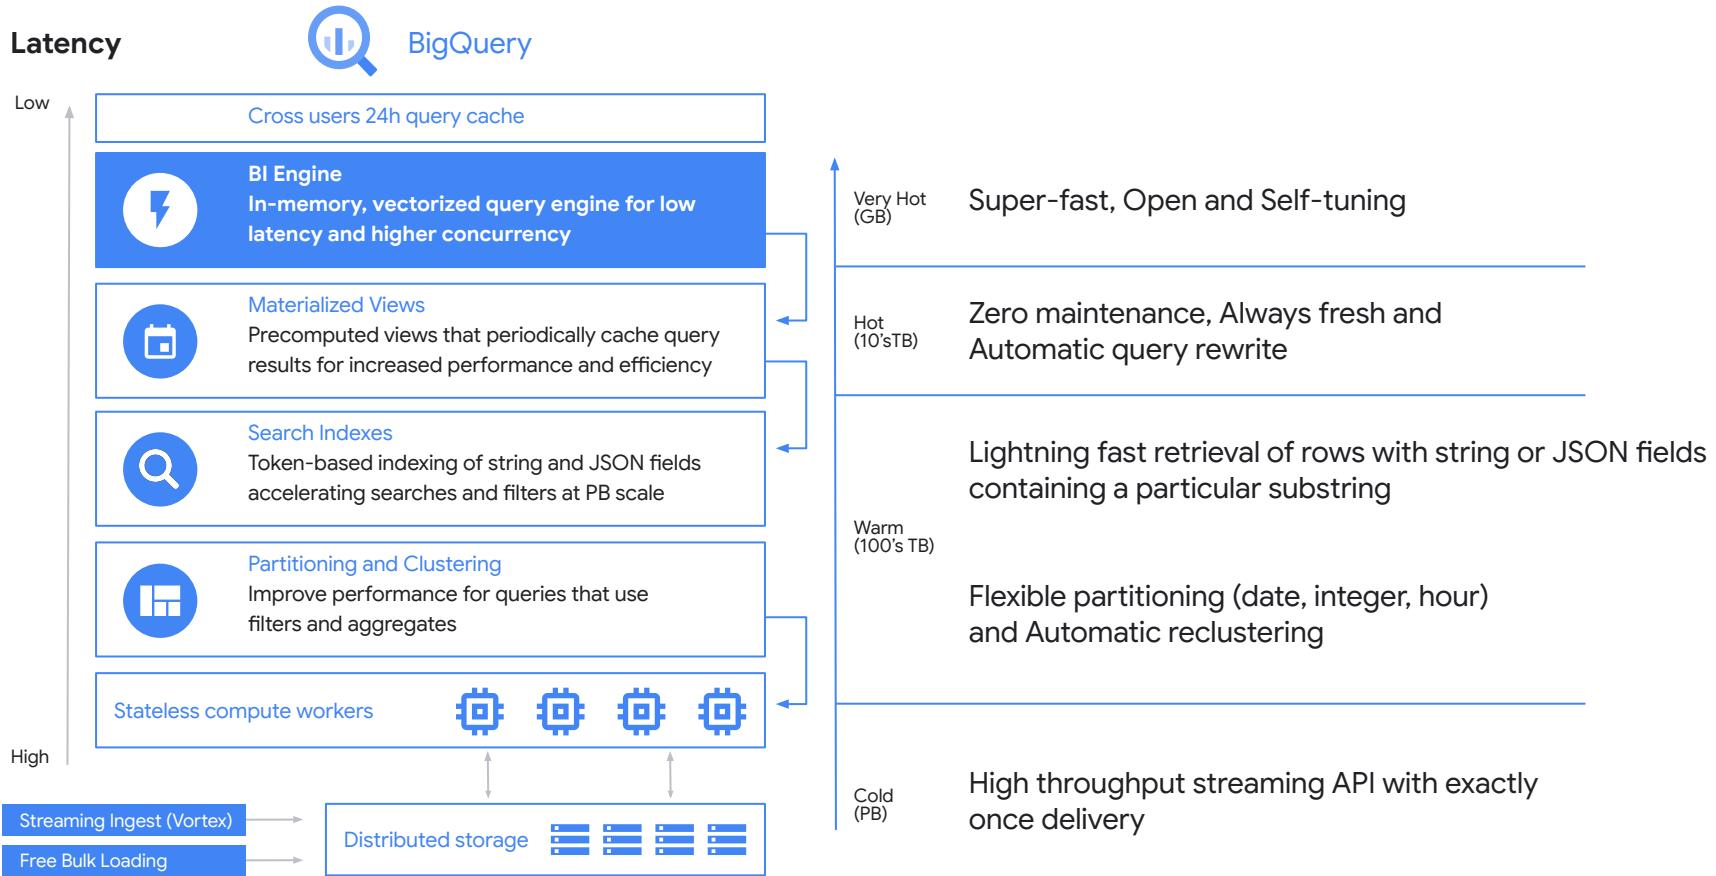
```
CREATE MATERIALIZED VIEW  project-id.my_dataset.my_mv_table
OPTIONS
(
    enable_refresh = true, refresh_interval_minutes = 60,
    max_staleness = INTERVAL "4" HOUR,
    allow_non_incremental_definition = true
)
AS SELECT
    employee_id,
    DATE(transaction_time),
    COUNT(1) AS count
FROM my_dataset.my_base_table
GROUP BY 1, 2;
```

Allows much wider SQL Support

N.B.

- Refresh will involve the entire materialized view.
- Can only directly query the materialized view, query rewrite will not work.

Adaptive Caching layers in BigQuery



BigQuery BI Engine

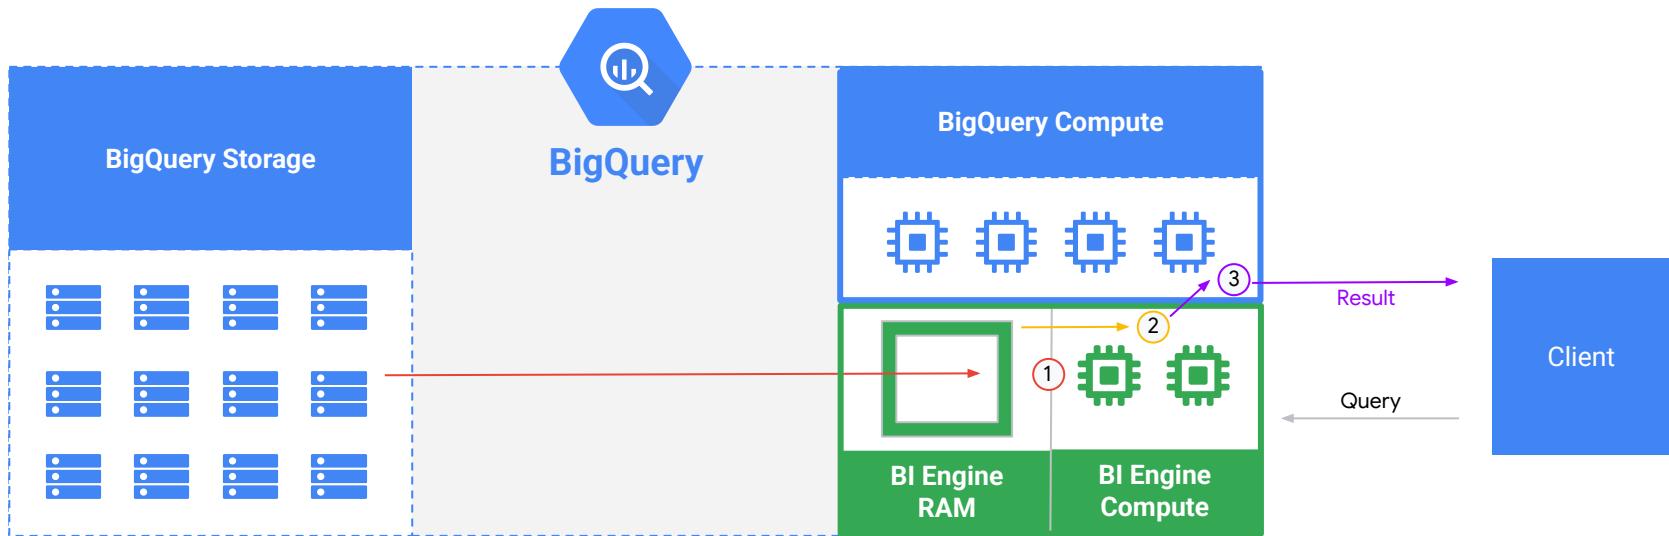
Sub-second queries

High Concurrency

Smart tuning



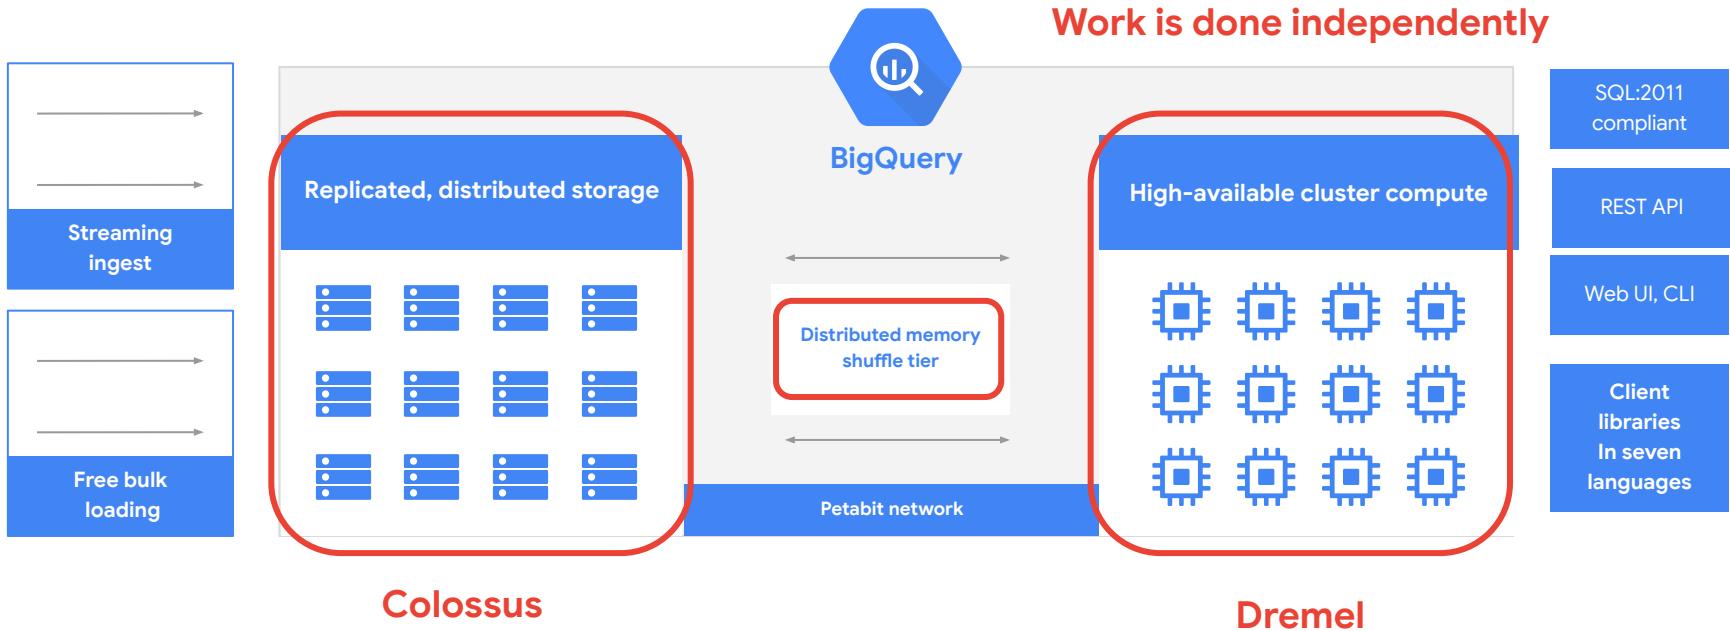
BI Engine



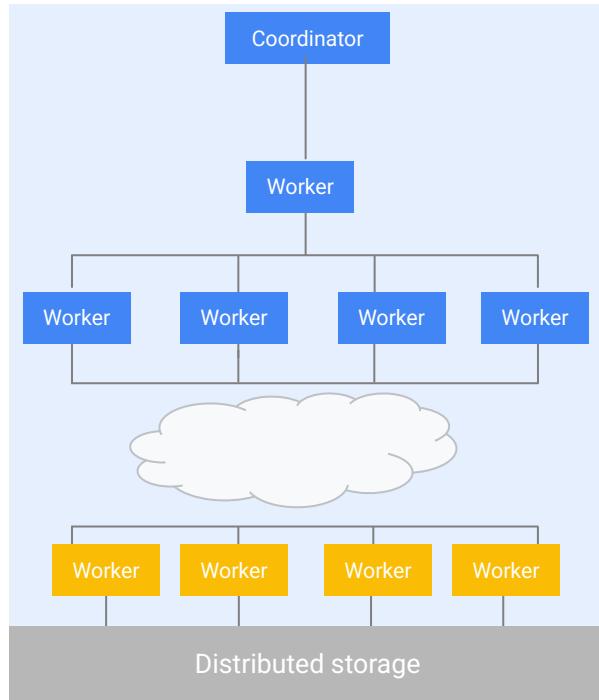
- ① Check if all data required to resolve query is in RAM. If not, scan additional from storage (bytes scanned are not billed).
- ② Use BI Engine compute to resolve leaf queries.
- ③ For complex queries, BQ compute may be used for final stages of query.



Going Deep - Query Processing Internals



Flexible Query Execution



```
SELECT language, MAX/views) as views  
FROM `wikipedia_benchmark.Wiki1B`  
WHERE title LIKE "G%o%"  
GROUP BY 1 ORDER BY 2 DESC LIMIT 100
```

Stage 3: SORT, LIMIT

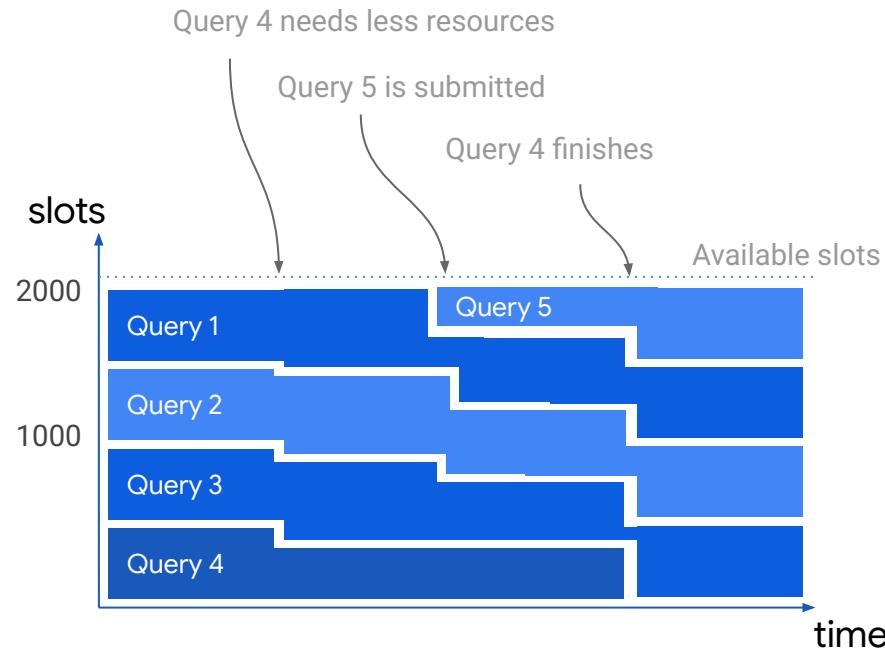
Stage 2: GROUP BY, SORT, LIMIT

Shuffle with dynamic partitioning

Stage 1: Partial GROUP BY

Dynamic Scheduling in BigQuery

- **Dynamic central scheduler allocates slots**
- **Handles machine failure**
- **Fair resource distribution between queries**



Processing location: US No cached results

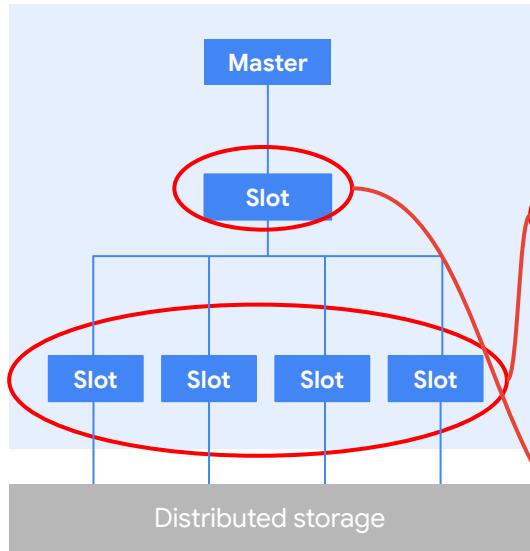
Run Save query Save view Schedule query More

Query results

Query complete (7.0 sec elapsed, 6.8 GB processed)

Job information Results JSON Execution details





Worker timing ⓘ

Stages	Wait	Read	Compute	Write	Input:	Rows
S00: Input	Avg: 54 ms Max: 111 ms	80 ms	54 ms	2 ms	Input: 58,937,715	19
READ	\$1:start_station_name FROM bigquery-public-data.new_york_citibike.citibike_trips WHERE like(\$1, '%Broadway%')				Output:	19
AGGREGATE	\$20 := COUNT_STAR()					
WRITE	\$20 TO __stage00_output					
S01: Output	Avg: 96 ms Max: 96 ms	0 ms	7 ms	3 ms	Input:	19
READ	\$20 FROM __stage00_output				Output:	1
AGGREGATE	\$10 := SUM_OF_COUNTS(\$20)					
WRITE	\$10 TO __stage01_output					



Focus on stages dominating resources



Pinpoint data skew (large diff in average and max time spent)



Reduce CPU tasks (consider approx functions)

```
job = client.get_job(  
    "bquxjob_7ae65493_179f5db3ab5"  
, location="US")
```

```
SELECT  
*  
FROM  
`region-us`.INFORMATION_SCHEMA.JOB_TIMELINE_BY_PROJECT  
WHERE  
job_id = "bquxjob_7ae65493_179f5db3ab5"
```



Going Deep -How to Optimise Queries

Optimization: Necessary columns only

Original code

```
SELECT  
*  
FROM  
`dataset.table`
```

Optimized

```
SELECT  
* EXCEPT (dim1, dim2)  
FROM  
`dataset.table`
```

Don't do "select * limit 10" to look at data, use the preview feature instead.

Optimization: Auto-pruning with Partitioning & Clustering

Table info

Table ID	bigquery-public-data:wikipedia.pageviews_2021
Table size	1.1 TB
Long-term storage size	497.28 GB
Number of rows	24,870,308,515
Created	Dec 31, 2020, 6:00:43 PM UTC-8
Last modified	Jun 11, 2021, 7:02:21 AM UTC-7
Table expiration	NEVER
Data location	US
Description	Wikipedia pageviews from http://dumps.wikimedia.you
Table Type	Partitioned
Partitioned by	DAY
Partitioned on field	datehour
Partition expiration	
Partition filter	Required
Clustered by	<input type="button" value="wiki"/> <input type="button" value="title"/>

Original code

```
SELECT
...
FROM
`wikipedia.pageviews_2021`
```

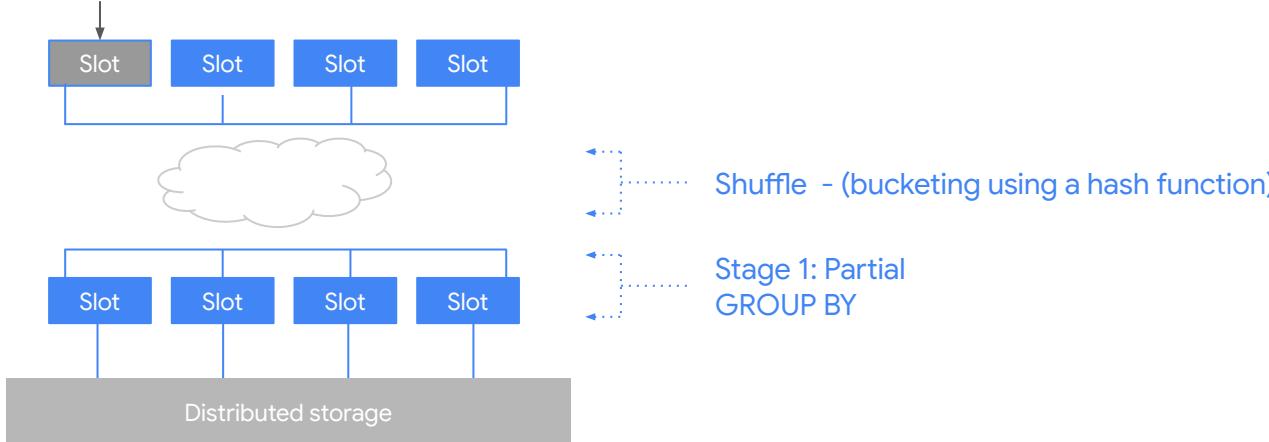
Optimized

```
SELECT
...
FROM
`wikipedia.pageviews_2021`
```

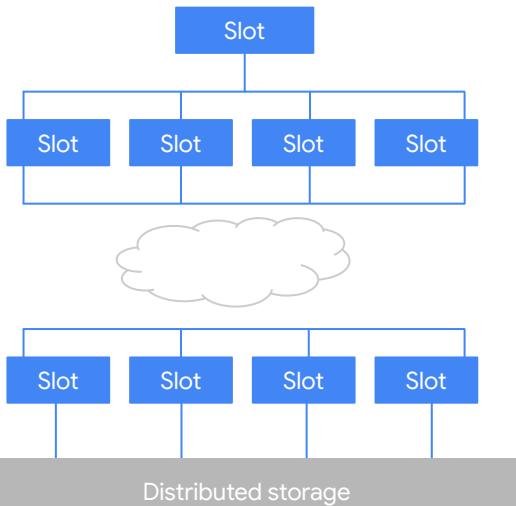
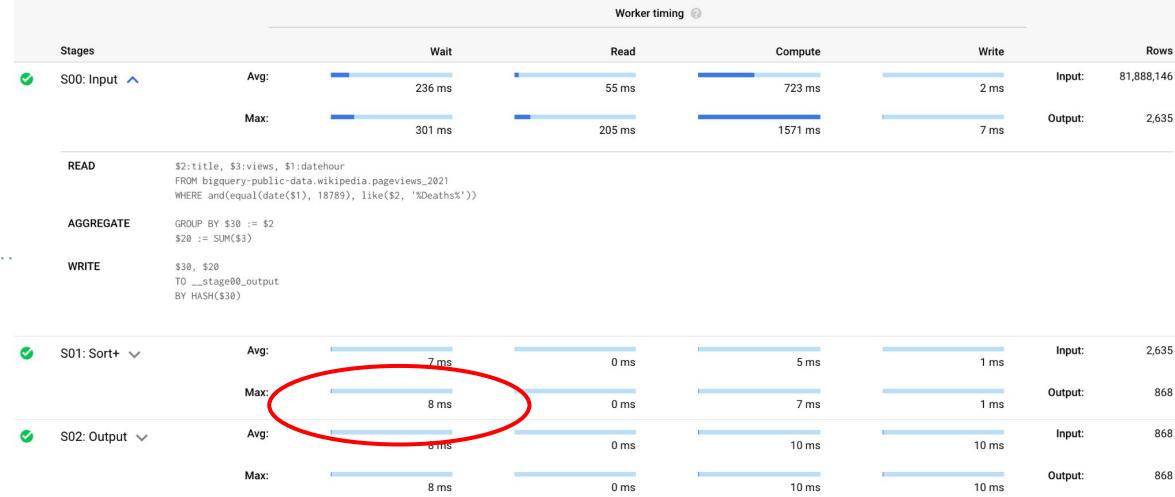
```
WHERE
DATE(datehour) = "2021-06-11"
AND wiki="simple"
```

```
SELECT title, SUM/views)
FROM `wikipedia.pageviews_2021`
WHERE DATE(datehour) = "2021-06-11"
GROUP BY title
ORDER BY SUM/views) DESC
LIMIT 1000
```

All partial aggregations for
the title “Games_in_2020”



```
SELECT title, SUM/views)
FROM `wikipedia.pageviews_2021`
WHERE DATE(datehour) = "2021-06-11"
GROUP BY title
ORDER BY SUM/views) DESC
LIMIT 1000
```



Optimization: Late aggregation

Original code

```
SELECT
    t1.dim1,
    SUM(t1.m1)
    SUM(t2.m2)
FROM (SELECT
    dim1,
    SUM(metric1) m1
    FROM `dataset.table1` GROUP BY 1) t1
JOIN (SELECT
    dim1,
    SUM(metric2) m2
    FROM `dataset.table2` GROUP BY 1) t2
ON t1.dim1 = t2.dim1
GROUP BY 1;
```

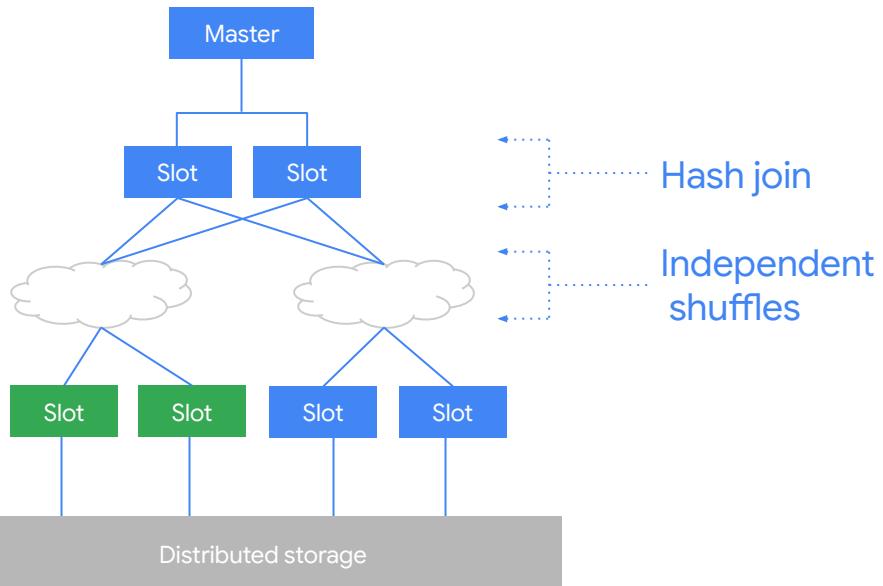
Optimized

```
SELECT
    t1.dim1,
    SUM(t1.m1)
    SUM(t2.m2)
FROM (SELECT
    dim1,
    metric1 m1
    FROM `dataset.table1` ) t1
JOIN (SELECT
    dim1,
    metric2 m2
    FROM `dataset.table2` ) t2
ON t1.dim1 = t2.dim1
GROUP BY 1;
```

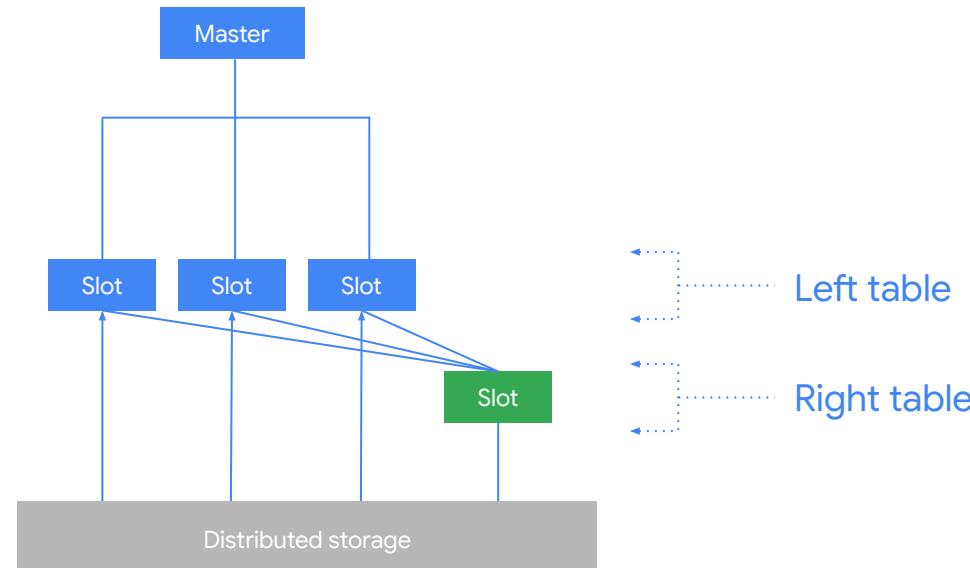
Avoid multiple group by executions, by moving them up.

Unless If a table can be reduced drastically by aggregating in preparation for being joined, then aggregate it early.

Large JOIN (shuffle)



Small JOIN (broadcast)



Optimization: JOIN pattern - largest table first

Original code

```
SELECT
    t1.dim1,
    SUM(t1.metric1),
    SUM(t2.metric2)
FROM
    `dataset.small_table` t1
JOIN
    `dataset.large_table` t2
ON
    t1.dim1 = t2.dim1
WHERE t1.dim1 = 'abc'
GROUP BY 1;
```

Optimized

```
SELECT
    t1.dim1,
    SUM(t1.metric1),
    SUM(t2.metric2)
FROM
    `dataset.large_table` t2
JOIN
    `dataset.small_table` t1
ON
    t1.dim1 = t2.dim1
WHERE t1.dim1 = 'abc'
GROUP BY 1;
```

Optimization: Filter before JOINs

Filter both tables before the join to reduce amount of data

Original code

```
SELECT
    t1.dim1,
    SUM(t1.metric1)
FROM
    `dataset.table1` t1
LEFT JOIN
    `dataset.table2` t2
ON
    t1.dim1 = t2.dim1
WHERE t2.dim2 = 'abc'
GROUP BY 1;
```

Optimized

```
SELECT
    t1.dim1,
    SUM(t1.metric1)
FROM
    `dataset.table1` t1
LEFT JOIN
    `dataset.table2` t2
ON
    t1.dim1 = t2.dim1
WHERE t2.dim2 = 'abc' AND t1.dim2 = 'abc'
GROUP BY 1;
```

WHERE clause: Expression order.

Original code

```
SELECT  
  text  
FROM  
  `stackoverflow.comments`  
WHERE  
  text LIKE '%java%'  
AND user_display_name = 'anon'
```

Optimized

```
SELECT  
  text  
FROM  
  `stackoverflow.comments`  
WHERE  
  user_display_name = 'anon'  
  AND text LIKE '%java%'
```

The first part of your where clause should always contain the filter that will eliminate the most data.

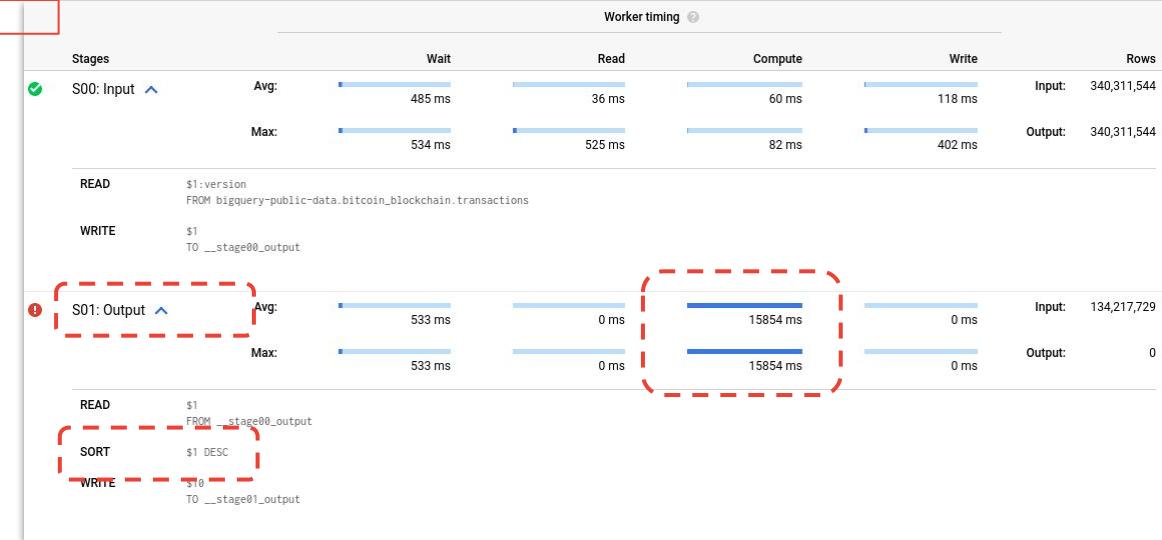
Optimization: ORDER BY with LIMIT

Original code

```
SELECT
    t.dim1,
    t.dim2,
    t.metric1
FROM
    `dataset.table` t
ORDER BY t.metric1 DESC
```

Optimized

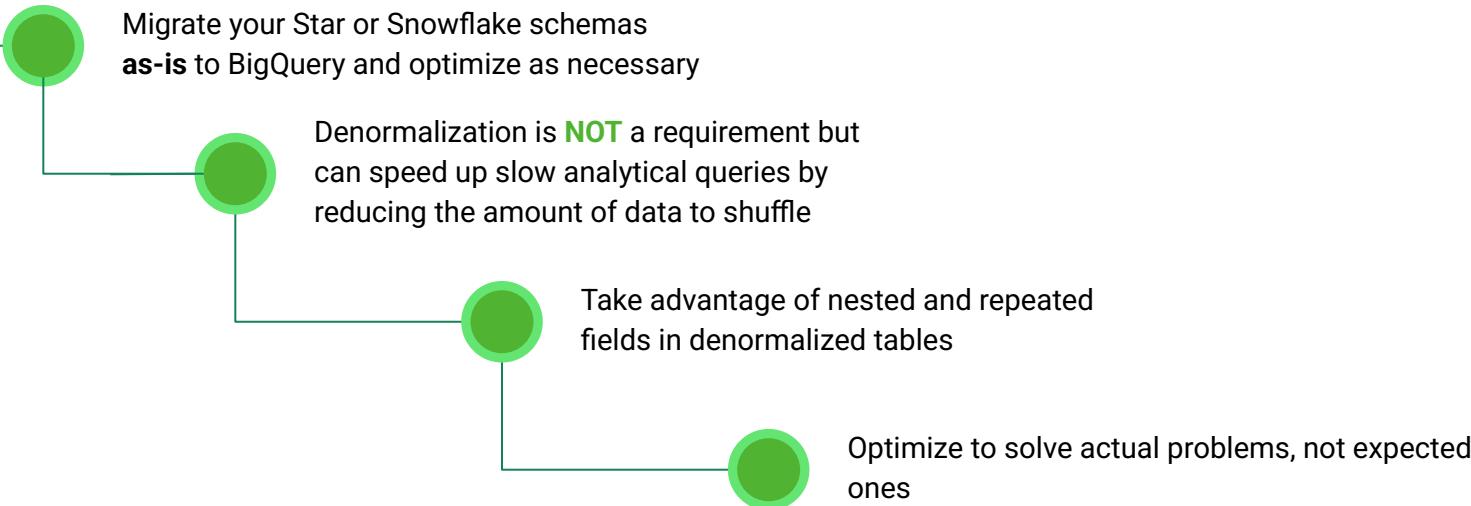
```
SELECT
    t.dim1,
    t.dim2,
    t.metric1
FROM
    `dataset.table` t
ORDER BY t.metric1 DESC
LIMIT 1000
```





Schema design and data organisation

Schema design in a nutshell



Nested and repeated columns

Orders		
Customer	Order_id	Date
Bob	1234	2020/01/02
Lisa	4321	2020/01/12
John	4567	2020/02/01

Order_items		
Order_id	Product	Quantity
1234	Jeans	2
1234	Shirt	4
4321	Skirt	1
4567	Boots	1
4567	Jeans	3
3567	Hat	1



Orders				
Customer	Order_id	Date	Item.Product	Item.Quantity
Bob	1234	2020/01/02	Jeans	2
			Shirt	4
Lisa	4321	2020/01/12	Skirt	1
John	4567	2020/02/01	Boots	1
			Jeans	3
			Hat	1

More natural representation of hierarchical data

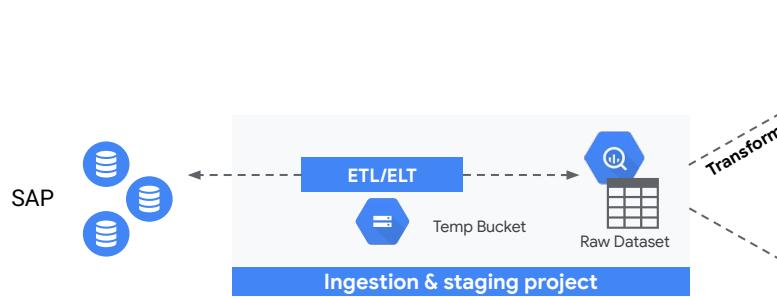
Can be accessed via dot notation:
name vs. Item[0].Quantity

Leverage Nested and Repeated fields for:

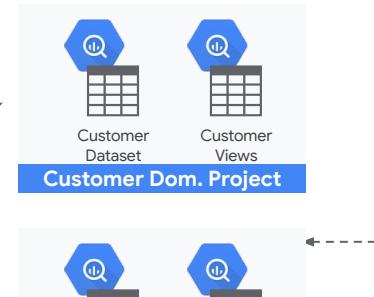
- Tightly-coupled or Immutable relationships
 - Session w/ Events
 - Order w/ Line Items
- Infrequently changing data
 - Country, Region, Date, etc

Store and organize your data on BigQuery (example)

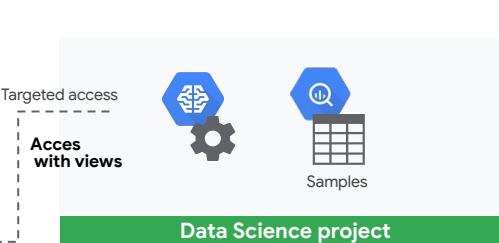
Raw Data Projects (source oriented)



Domain Data Projects (data as a product oriented)



Usage Data Projects (data as a product oriented)



Composer



DataPlex



Monitoring



Logging

Shared tools project

Thank You

Summary

- Storage
 - Partitioning
 - Clustering
 - Search indexes
 - Materialized views
 - Bi Engine
- Compute
 - Slots
 - Query execution
- SQL best practices
- Q&A



Q&A